


Network Programming with Sockets

CS 475

Anatomy of an Internet Connection



- 2 -

Sockets

What is a socket?


- To the kernel, a socket is an endpoint of communication.
- To an application, a socket is a file descriptor that lets the application read/write from/to the network.
 - Remember: All Unix I/O devices, including networks, are modeled as files.

Clients and servers communicate with each other by reading from and writing to socket descriptors.

The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors.

– 3 –

Overview of the Sockets Interface




– 4 –

Socket Address Structures

Generic socket address:

- For address arguments to connect, bind, and accept.
- Necessary only because C did not have generic (`void *`) pointers when the sockets interface was designed.

```
struct sockaddr {  
    unsigned short  sa_family;    /* protocol family */  
    char           sa_data[14];   /* address data. */  
};
```




- 5 -

Socket Address Structures

Internet-specific socket address:

- Must cast (`sockaddr_in *`) to (`sockaddr *`) for connect, bind, and accept.

```
struct sockaddr_in {  
    unsigned short  sin_family;    /* address family (always AF_INET) */  
    unsigned short  sin_port;     /* port num in network byte order */  
    struct in_addr  sin_addr;    /* IP addr in network byte order */  
    unsigned char   sin_zero[8]; /* pad to sizeof(struct sockaddr) */  
};
```



- 6 -

Example: Echo Client and Server

On Server

```
bass> echoserver 5000
server established connection with KITTYHAWK.CMCL (128.2.194.242)
server received 4 bytes: 123
server established connection with KITTYHAWK.CMCL (128.2.194.242)
server received 7 bytes: 456789
...
```

On Client

```
kittyhawk> echoclient bass 5000
Please enter msg: 123
Echo from server: 123

kittyhawk> echoclient bass 5000
Please enter msg: 456789
Echo from server: 456789
kittyhawk>
```

- 7 -

Echo Client Main Routine

```
#include "csapp.h"


/* usage: ./echoclient host port */
int main(int argc, char **argv)
{
    int clientfd, port;
    char *host, buf[MAXLINE];
    rio_t rio;
    host = argv[1];  port = atoi(argv[2]);
    clientfd = Open_clientfd(host, port);
    Rio_readinitb(&rio, clientfd);
    printf("type:");
    ffflush(stdout);
    while (Fgets(buf, MAXLINE, stdin) != NULL) {
        Rio_writen(clientfd, buf, strlen(buf));
        Rio_readlineb(&rio, buf, MAXLINE);
        printf("echo:");
        Fputs(buf, stdout);
        printf("type:");
        ffflush(stdout);
    }
    Close(clientfd);
    exit(0);
}
```

Send line to server

Receive line from server

- 8 -

Overview of the Sockets Interface



- 9 -

Echo Client: `open_clientfd`

```
int open_clientfd(char *hostname, int port) {  
    int clientfd;  
    struct hostent *hp;  
    struct sockaddr_in serveraddr;  
  
    if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)  
        return -1; /* check errno for cause of error */  
  
    /* Fill in the server's IP address and port */  
    if ((hp = gethostbyname(hostname)) == NULL)  
        return -2; /* check h_errno for cause of error */  
    bzero((char *) &serveraddr, sizeof(serveraddr));  
    serveraddr.sin_family = AF_INET;  
    bcopy((char *)hp->h_addr_list[0],  
          (char *)&serveraddr.sin_addr.s_addr, hp->h_length);  
    serveraddr.sin_port = htons(port);  
  
    /* Establish a connection with the server */  
    if (connect(clientfd, (SA *) &serveraddr,  
               sizeof(serveraddr)) < 0)  
        return -1;  
    return clientfd;  
}
```

This function opens a connection from the client to the server at `hostname:port`

>Create socket

Create address

Establish connection

Echo Client: open_clientfd (socket)

socket creates a socket descriptor on the client

- Just allocates & initializes some internal data structures
- **AF_INET**: indicates that the socket is associated with Internet protocols.
- **SOCK_STREAM**: selects a reliable byte stream connection
 - Provided by TCP

```
int clientfd; /* socket descriptor */

if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
    return -1; /* check errno for cause of error */

... (more)
```

- 11 -

Echo Client: open_clientfd (gethostbyname)

The client then builds the server's Internet address.

```
int clientfd; /* socket descriptor */
struct hostent *hp; /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */

...

/* fill in the server's IP address and port */
if ((hp = gethostbyname(hostname)) == NULL)
    return -2; /* check h_errno for cause of error */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_port = htons(port);
bcopy((char *)hp->h_addr_list[0],
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
```

Check this out!

- 12 -

A Careful Look at bcopy Arguments

```
/* DNS host entry structure */
struct hostent {
    ...
    int    h_length;      /* length of an address, in bytes */
    char   **h_addr_list; /* null-terminated array of in_addr structs */
};

struct sockaddr_in {
    ...
    struct in_addr  sin_addr;    /* IP addr in network byte order */
    ...
}; /* Internet address structure */
struct in_addr {
    unsigned int s_addr; /* network byte order (big-endian) */
};

struct hostent *hp;           /* DNS host entry */
struct sockaddr_in serveraddr; /* server's IP address */
...
bcopy((char *)hp->h_addr_list[0], /* src, dest */
      (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
```

- 13 -

Echo Client: open_clientfd (connect)

Finally the client creates a connection with the server.

- Client process suspends (blocks) until the connection is created.
- After resuming, the client is ready to begin exchanging messages with the server via Unix I/O calls on descriptor clientfd.

```
int clientfd;           /* socket descriptor */
struct sockaddr_in serveraddr; /* server address */
typedef struct sockaddr SA; /* generic sockaddr */
...
/* Establish a connection with the server */
if (connect(clientfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;
return clientfd;
}
```

- 14 -

Echo Server: Main Routine


```
int main(int argc, char **argv) {
    int listenfd, connfd, port, clientlen;
    struct sockaddr_in clientaddr;
    struct hostent *hp;
    char *haddrp;

    port = atoi(argv[1]); /* the server listens on a port passed
                           on the command line */
    listenfd = open_listenfd(port);

    while (1) {
        clientlen = sizeof(clientaddr);
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                           sizeof(clientaddr.sin_addr.s_addr), AF_INET);
        haddrp = inet_ntoa(clientaddr.sin_addr);
        printf("server connected to %s (%s)\n", hp->h_name, haddrp);
        echo(connfd);
        Close(connfd);
    }
}
```

- 15 -

Overview of the Sockets Interface



- 16 -

Echo Server: open_listenfd

```
int open_listenfd(int port)
{
    int listenfd, optval=1;
    struct sockaddr_in serveraddr;

    /* Create a socket descriptor */
    if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1;

    /* Eliminates "Address already in use" error from bind. */
    if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
                   (const void *)&optval , sizeof(int)) < 0)
        return -1;

    ... (more)
```

- 17 -

Echo Server: open_listenfd (cont)

```
...

/* Listenfd will be an endpoint for all requests to port
   on any IP address for this host */
bzero((char *) &serveraddr, sizeof(serveraddr));
serveraddr.sin_family = AF_INET;
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
serveraddr.sin_port = htons((unsigned short)port);
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
    return -1;

/* Make it a listening socket ready to accept
   connection requests */
if (listen(listenfd, LISTENQ) < 0)
    return -1;

return listenfd;
}
```

- 18 -

Echo Server: `open_listenfd` (socket)

socket creates a socket descriptor on the server.

- `AF_INET`: indicates that the socket is associated with Internet protocols.
- `SOCK_STREAM`: selects a reliable byte stream connection (TCP)

```
int listenfd; /* listening socket descriptor */  
  
/* Create a socket descriptor */  
if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)  
    return -1;
```

- 19 -

Echo Server: `open_listenfd` (setsockopt)

The socket can be given some attributes.

```
...  
/* Eliminates "Address already in use" error from bind(). */  
if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,  
               (const void *)&optval , sizeof(int)) < 0)  
    return -1;
```

Handy trick that allows us to rerun the server immediately after we kill it.

- Otherwise we would have to wait about 15 secs.
- Eliminates “Address already in use” error from `bind()` .

Strongly suggest you do this for all your servers to simplify debugging.


- 20 -

Echo Server: open_listenfd (initialize socket address)

Initialize socket with server port number

accept connection from any IP address

```
struct sockaddr_in serveraddr; /* server's socket addr */  
...  
/* listenfd will be an endpoint for all requests to port  
on any IP address for this host */  
bzero((char *) &serveraddr, sizeof(serveraddr));  
serveraddr.sin_family = AF_INET;  
serveraddr.sin_port = htons((unsigned short)port);  
serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
```



IP addr and port stored in network (big-endian) byte order

- 21 -

Echo Server: open_listenfd (bind)

bind associates the socket with the socket address we just created.

```
int listenfd; /* listening socket */  
struct sockaddr_in serveraddr; /* server's socket addr */  
...  
/* listenfd will be an endpoint for all requests to port  
on any IP address for this host */  
if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)  
    return -1;
```

- 22 -

Echo Server: `open_listenfd` (`listen`)

`listen` indicates that this socket will accept connection (`connect`) requests from clients

`LISTENQ` is constant indicating how many pending requests allowed

```
int listenfd; /* listening socket */  
  
...  
/* Make it a listening socket ready to accept connection requests */  
if (listen(listenfd, LISTENQ) < 0)  
    return -1;  
return listenfd;  
}
```

We're finally ready to enter the main server loop that accepts and processes client connection requests.

- 23 -


Echo Server: Main Loop

The server loops endlessly, waiting for connection requests, then reading input from the client, and echoing the input back to the client.

```
main() {  
  
    /* create and configure the listening socket */  
  
    while(1) {  
        /* Accept(): wait for a connection request */  
        /* echo(): read and echo input lines from client til EOF */  
        /* Close(): close the connection */  
    }  
}
```

- 24 -

Overview of the Sockets Interface



Echo Server: accept

`accept()` blocks waiting for a connection request.

```
int listenfd; /* listening descriptor */
int connfd; /* connected descriptor */
struct sockaddr_in clientaddr;
int clientlen;

clientlen = sizeof(clientaddr);
connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
```


`accept` returns a **connected descriptor** (`connfd`) with the same properties as the **listening descriptor** (`listenfd`)

- Returns when the connection between client and server is created and ready for I/O transfers.
- All I/O with the client will be done via the connected socket.

`accept` also fills in client's IP address.

- 26 -

Echo Server: accept Illustrated



- 27 -

Connected vs. Listening Descriptors

Listening descriptor

- End point for client connection requests.
- Created once and exists for lifetime of the server.

Connected descriptor

- End point of the connection between client and server.
- A new descriptor is created each time the server accepts a connection request from a client.
- Exists only as long as it takes to service client.

Why the distinction?

- Allows for concurrent servers that can communicate over many client connections simultaneously.
 - E.g., Each time we receive a new request, we fork a child to handle the request.

- 28 -

Echo Server: Identifying the Client

The server can determine the domain name and IP address of the client.

```
struct hostent *hp; /* pointer to DNS host entry */
char *haddrp; /* pointer to dotted decimal string */

hp = gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                    sizeof(clientaddr.sin_addr.s_addr), AF_INET);
haddrp = inet_ntoa(clientaddr.sin_addr);
printf("server connected to %s (%s)\n", hp->h_name, haddrp);
```

- 29 -

Echo Server: echo

The server uses RIO to read and echo text lines until EOF (end-of-file) is encountered.

- EOF notification caused by client calling `close(clientfd)` .
- **IMPORTANT:** EOF is a condition, not a particular data byte.

```
void echo(int connfd)
{
    size_t n;
    char buf[MAXLINE];
    rio_t rio;

    Rio_readinitb(&rio, connfd);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
        upper_case(buf);
        Rio_writen(connfd, buf, n);
        printf("server received %d bytes\n", n);
    }
}
```

The RIO Package

RIO is a set of wrappers that provide efficient and robust I/O in apps, such as network programs that are subject to short counts

RIO provides two different kinds of functions

- Unbuffered input and output of binary data
 - `rio_readn` and `rio_writen`
- Buffered input of binary data and text lines
 - `rio_readlineb` and `rio_readnb`
 - Buffered RIO routines are *thread-safe* and can be interleaved arbitrarily on the same descriptor

Included in file `csapp.c`, `csapp.h` provided with Assignment 3

- 31 -

Unbuffered RIO Input and Output

Same interface as Unix `read` and `write`

Especially useful for transferring data on network sockets

```
#include "csapp.h"

ssize_t rio_readn(int fd, void *usrbuf, size_t n);
ssize_t rio_writen(int fd, void *usrbuf, size_t n);

Return: num. bytes transferred if OK, 0 on EOF (rio_readn only), -1 on error
```

- `rio_readn` returns short count only if it encounters EOF.
 - Only use it when you know how many bytes to read
- `rio_writen` never returns a short count.
- Calls to `rio_readn` and `rio_writen` can be interleaved arbitrarily on the same descriptor.

- 32 -

RIO Example

Copying the lines of a text file from standard input to standard output

```
#include "csapp.h"

int main(int argc, char **argv)
{
    int n;
    rio_t rio;
    char buf[MAXLINE];

    Rio_readinitb(&rio, STDIN_FILENO);
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0)
        Rio_writen(STDOUT_FILENO, buf, n);
    exit(0);
}
```

- 33 -

Standard I/O: Reading Files

Reading a file copies bytes from the current file position to memory, and then updates file position

```
char buf[512];
int fd;          /* file descriptor */
int nbytes;     /* number of bytes read */

/* Open file fd ... */
/* Then read up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

Returns number of bytes read from file `fd` into `buf`

- Return type `ssize_t` is signed integer
- `nbytes < 0` indicates that an error occurred.
- **short counts** (`nbytes < sizeof(buf)`) are possible and are not errors!

- 34 -

Motivation for RIO: Dealing with Short Counts

Short counts can occur in these situations:

- Encountering (end-of-file) EOF on reads
- Reading text lines from a terminal
- Reading and writing network sockets or Unix pipes

Short counts never occur in these situations:

- Reading from disk files (except for EOF)
- Writing to disk files

One way to deal with short counts in your code:

- Use the RIO (Robust I/O) package

– 35 –

Testing Servers Using telnet

The telnet program is invaluable for testing servers that transmit ASCII strings over Internet connections

- Our simple echo server
- Web servers
- Mail servers

Usage:

- `unix> telnet <host> <portnumber>`
- Creates a connection with a server running on `<host>` and listening on port `<portnumber>`.

– 36 –

Testing the Echo Server With telnet

```
bass> echoserver 5000
server established connection with KITTYHAWK.CMCL (128.2.194.242)
server received 5 bytes: 123
server established connection with KITTYHAWK.CMCL (128.2.194.242)
server received 8 bytes: 456789

kittyhawk> telnet bass 5000
Trying 128.2.222.85...
Connected to BASS.CMCL.CS.CMU.EDU.
Escape character is '^].
123
123
Connection closed by foreign host.
kittyhawk> telnet bass 5000
Trying 128.2.222.85...
Connected to BASS.CMCL.CS.CMU.EDU.
Escape character is '^].
456789
456789
Connection closed by foreign host.
kittyhawk>
```

- 37 -

For More Information

W. Richard Stevens, “Unix Network Programming: Networking APIs: Sockets and XTI”, Volume 1, Second Edition, Prentice Hall, 1998.

- THE network programming bible.

Unix Man Pages

- Good for detailed information about specific functions

Complete versions of the echo client and server are developed in the CS 367 text (Bryant & O'Halloran)

- Available from csapp.cs.cmu.edu
- You should compile and run them for yourselves to see how they work.
- Feel free to borrow any of this code.

- 38 -